

Random Forests and Decision Trees

```
=====
author: An alternative to predicting credit defaulters -
which model performs better?
date:
autosize: true
```

Previously

```
=====
- Previously we used logistic regression to predict whether
or not someone was going to default on credit payment
- However, logistic regression is not the only model
available
- In this section we will cover random forests and decision
trees, an alternative to linear and logistic regression. We
will compare models and decide which is more suited to our
needs.
```

What is a tree?

```
=====
- Given a dataset, we split the data recursively so that
each subset is as homogenous as possible
- We measure homogeneity using the Gini index. Alternatively
we can also use something called cross-entropy
- Both the Gini Index and cross-entropy measure the total
variance across all classes in our data.
- A small value means out split is "purer"
```

Visualising Splits

```
=====
![alt text](TreeSplits.png)
```

Source: Introduction to Statistical Learning

Pruning a tree

```
=====
- The problem: too many splits lead to us overfitting on our
data, too few and we underfit
- Usually the solution is to grow a large tree then prune it
back to a smaller subtree
```

Let's use Decision Trees on our Default Dataset

```
=====
Perform train/test split
```{r}
library(MASS)
library(ISLR)
library(caret)
```

```
set.seed(88)
train_Index <- createDataPartition(Default$default, p=0.8,
list=FALSE)
train_credit <- Default[train_Index,]
test_credit <- Default[-train_Index,]
head(train_credit)
```
```

Fit Model

```
=====
```{r}
modell <- train(default ~ ., data=train_credit,
method='rpart')
plot(modell)
by default, our complexity parameter is 3
```
```

Playing with different values for the complexity parameter

```
=====
```{r}
model2 <- train(default ~ balance, data=train_credit,
method='rpart', tuneLength=6)
plot(model2)
```
```

```
=====
```{r}
model2
```
```

Evaluating our model with test data

```
=====
```{r}
pred <- predict(model2, test_credit)
confusionMatrix(pred, test_credit$default)
```
```

ROC Curve

```
=====
```{r}
library(pROC)
pred_prob <- predict(model2, test_credit, type="prob")
rpart_ROC <- roc(predictor = pred_prob$No, response =
test_credit$default, levels =
rev(levels(test_credit$default)))
plot(rpart_ROC)
```
```

The problem with a single tree

```
=====
- Decision Trees suffer from high variance
- Fitting our model to distinct data sets could lead to quite different results.
- In contrast, models such as linear regression tend to have low variance.
- To overcome high variance, we use methods such as Random Forests.
- Essentially, we take repeated samples from one training dataset, fit a model to each sample, then average the predictions. The idea is that averaging in this way reduces variance.
```

Random Forests

```
=====
```{r}
model_rf <- train(default ~ student + balance + income,
data=Default, method="rf")

pred = predict(model_rf, newdata=test_credit)
confusionMatrix(data=pred, test_credit$default)
```
```

Comparing Decision Trees and Random Forests to Logistic Regression and Linear Regression

```
=====
- Classical approaches like linear regression assume a linear relationship between dependent and independent variables.
- If this assumption is true, a linear model will likely outperform a tree model
- However, if the relationship between variables is complex and non-linear, then a tree model might perform better.
- In terms of interpretability, trees are easier to visualise and understand
- Trees can handle qualitative predictors without having to create dummy variables.
```

![alt text](TreesandLinearModels.png)

Introduction to K-fold Cross Validation

```
=====
When evaluating models, we often want to assess how well it performs in predicting the target variable on different subsets of the data. One such technique for doing this is k-
```

fold cross-validation, which partitions the data into k equally sized segments (called 'folds'). One fold is held out for validation while the other k-1 folds are used to train the model and then used to predict the target variable in our testing data. This process is repeated k times, with the performance of each model in predicting the hold-out set being tracked using a performance metric such as accuracy. The most common variation of cross validation is 10-fold cross-validation.

```
```{r}
implement K-fold Cross Validation in caret
ctrl <- trainControl(method = "repeatedcv", number = 10,
savePredictions = TRUE)
```

```
model_with_cv <- train(default ~ student + balance + income,
data=Default, method="rf",
trControl = ctrl)
```

```
pred = predict(model_with_cv, newdata=test_credit)
confusionMatrix(data=pred, test_credit$default)
```
```